

Práctica 1: Sampling and Quantization

Renato Bedriñana Cárdenas

Hugo Blanco Demelo

23 de octubre de 2025

1 Task 1

Question: Give your interpretation of the resulting graphs. Do the quantization levels correspond with the values you had expected?


We represent the original continuous signal x (blue) and the 2 quantized version using 2 (red) and 4 (yellow) bits. As expected, the 2 bit quantization produces fewer discrete levels than the 4 bits quantization. Increase the number of bits decreases Δ , resulting in smaller steps and a quantized signal that follows the input more closely.

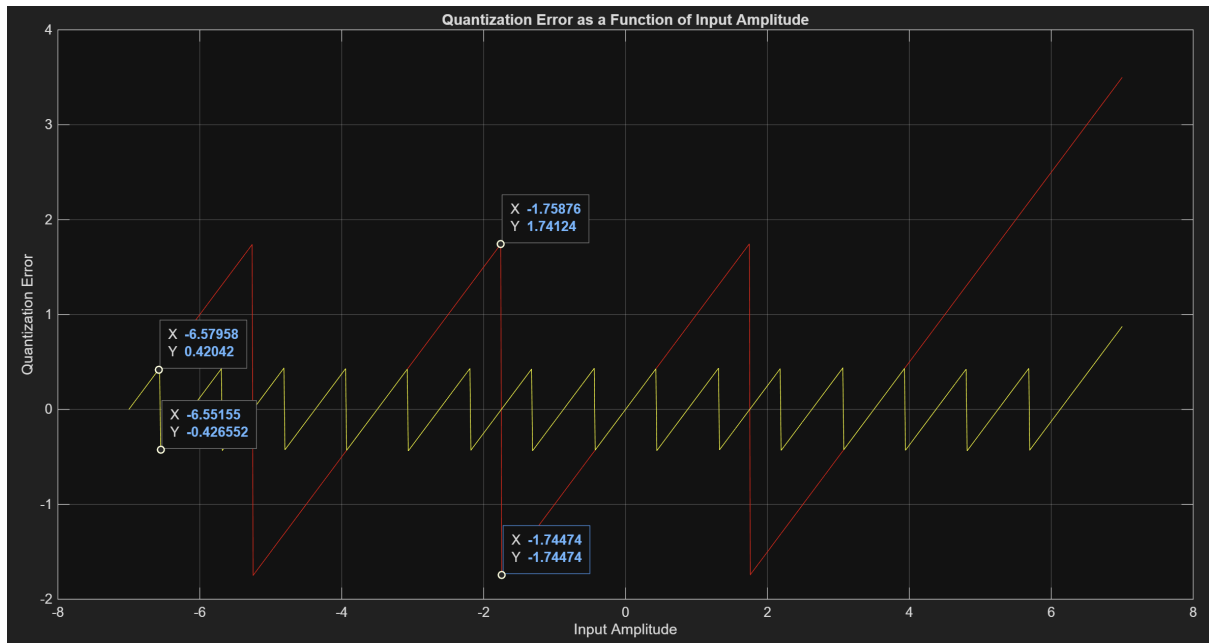
Question: For both cases, represent the quantization error as a function of input amplitude in the range $[-7, +7]$ and comment on your results. Is this error always within the $[-\Delta/2, +\Delta/2]$ interval?

The magnitude of the error decreases as the number of bits increases, since a smaller quantization step Δ reduces the maximum deviation between the input and its quantized version. The $[-\Delta/2, +\Delta/2]$ in each case is as follows:

- For $N = 2$ the Δ value we get is $\Delta = 3,5$, so the interval should be $[-1,75, 1,75]$.
- For $N = 4$ the Δ value we get is $\Delta = 0,875$, so the interval should be $[-0,4375, 0,4375]$.

In both cases, the error remains bounded within the theoretical interval $[-\Delta/2, +\Delta/2]$.






2 Task 2

Question: Assume a full-scale sinusoidal input and plot the histogram of the quantization error. Do you observe what you expected, or not?

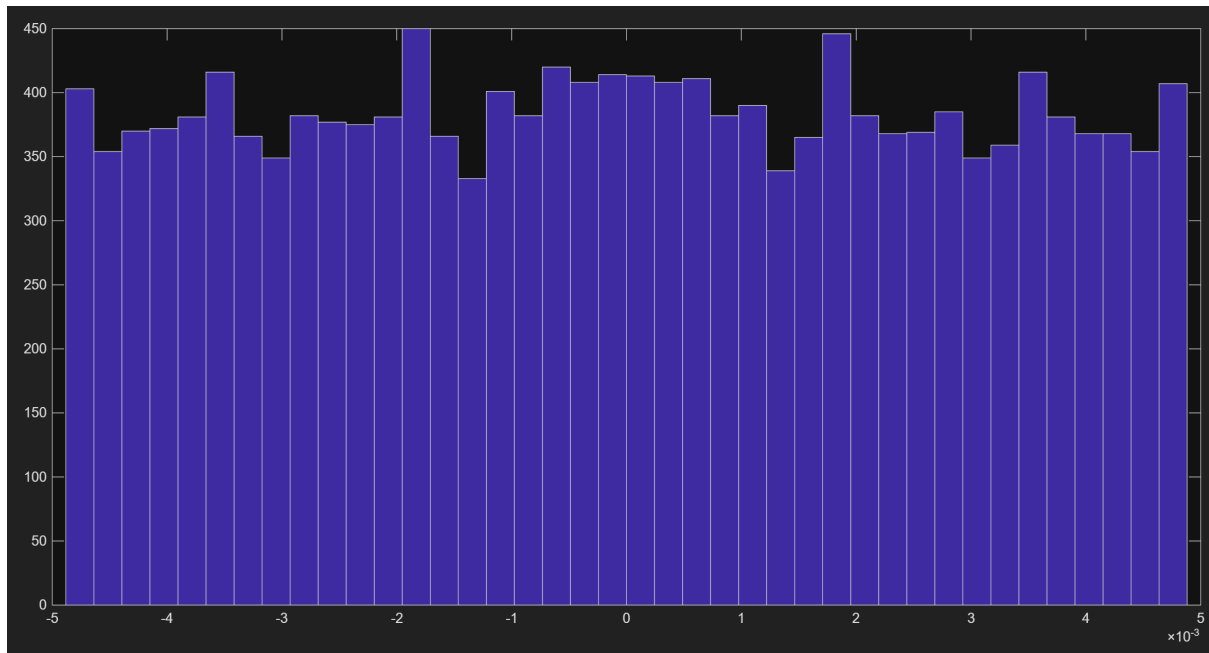
$\Delta = \frac{2*FS}{2^N} = 0,0098$, so the $[-\frac{\Delta}{2}, +\frac{\Delta}{2}]$ interval should be $[-0,0049, +0,0049]$. In the histogram we can see that in that interval the error is uniformly distributed, but there is an error tail in the positive extreme. It means that there is **clipping** in the positive.



Question: Explain the operation of the Matlab command `var`. Estimate the variance of the quantization error using `var`, and compare it to its theoretical value. Estimate the value (in dB) of the Signal-to-Quantization Noise Ratio (SQNR) and compare it to its theoretical value (1).

The MATLAB command `var` computes the variance of a set of values. For a vector x , it calculates: $\text{var}(x) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$, where \bar{x} is the mean of the values in x .

ADD



Question: Repeat the previous steps for sinusoids with different amplitudes, and with decreasing resolutions of 12, 10, 8, 6 and 4 bits, in order to fill Table 1, rounding the SQNR values (in dB) to two decimal places. Comment on your results.

	$A = 0,5 \cdot \text{FS}$		$A = 0,75 \cdot \text{FS}$		$A = \text{FS}$		$A = 1,03 \cdot \text{FS}$	
	SQNR (dB)		SQNR (dB)		SQNR (dB)		SQNR (dB)	
N	theory	measured	theory	measured	theory	measured	theory	measured
12	61.96	62.02	67.98	68.03	71.50	71.54	74.00	73.70
10	49.92	50.01	55.94	56.01	59.46	59.53	61.96	61.56
8	37.88	38.07	43.90	44.04	47.42	47.53	49.92	49.15
6	25.84	26.22	31.86	32.13	35.38	35.60	37.88	36.52
4	13.80	14.59	19.82	20.37	23.34	23.78	25.84	23.63

Cuadro 1: Pertaining to Task 2.

3 Task 3

Question: Suppose that you have an N -bit A/D converter with tunable FS, and you know that your input samples follow a symmetric triangular pdf in some interval

$[-x_0, x_0]$. Intuitively, how would you set the FS value of your converter? What would the resulting rms value σ_x in dBFS be?

If you set $FS < x_0$ any input $|x|$ greater than FS will be clipped. If $FS > x_0$, we would be wasting the converter's since the signal would never reach the limits. Therefore, the value of FS should be x_0 .

$$\sigma_x = \sqrt{\text{var}(x)} = \frac{x_0}{\sqrt{6}} \text{ and in dBFS would be } 20 \log_{10}(1/\sqrt{6}) = -7.78 \text{ dBFS.}$$

Question: Explain how to generate in Matlab samples of a random variable following a symmetric triangular pdf with zero mean and rms value σ_0 . Check the histogram and use the commands mean and var to validate your approach

We have two options to do it:

- Option 1: We can do it using *makedist* function. To do it, we can use the following code:

```
x0 = 2;
A = -x0; B = 0; C = +x0; % simetria = media 0

pd = makedist('Triangular','A',A,'B',B,'C',C);
N = 100000;
samples = random(pd, N, 1);

% comprobaciones rapidas
emp_mean = mean(samples);
emp_var = var(samples);
emp_desv_std = std(samples);

% valores teoricos
% theo_mean = 0; % simetria centrado en 0
theo_var = (A^2 + B^2 + C^2 - A*B - A*C - B*C)/18;
rms = 20*log10(sqrt(theo_var)/x0);

fprintf('Theorical mean: 0; emp mean: %.2f\n',emp_mean)
;
fprintf('Theorical var: %.2f; emp var: %.2f\n',theo_var
,emp_var);
fprintf('Sigma value: %.2f\n',sqrt(theo_var));
fprintf('rms value in dBFS: %.2f\n',rms)

% ver histograma y pdf teorica
xgrid = linspace(A,C,500)';
figure
histogram(samples,100,'Normalization','pdf')
hold on
plot(xgrid, pdf(pd,xgrid), 'LineWidth',1.5)
title('Triangular (media 0) -- muestras vs PDF')
hold off
```

- Option 2: we can generate samples of a random variable following a symmetric triangular pdf as the sum of two independent random variables X_1 and X_2 from a uniform distribution. we can do it as follows: REVISAR!!

```
x0=2;
```

```

sigma0 = x0/sqrt(2);
N = 100000;

c = sigma0 * sqrt(3/2);

x1 = (2 * rand(N, 1) - 1) * c;
x2 = (2 * rand(N, 1) - 1) * c;

y = x1 + x2;

sample_mean = mean(y);
sample_var = var(y);
sample_rms = std(y);

fprintf('--- Validation ---\n');
fprintf('Target Mean: 0.0\n');
fprintf('Sample Mean: %f\n\n', sample_mean);

fprintf('Target Variance (sigma0^2): %f\n', sigma0^2);
fprintf('Sample Variance: %f\n\n', sample_var);

fprintf('Target RMS (sigma0): %f\n', sigma0);
fprintf('Sample RMS: %f\n\n', sample_rms);

figure;
histogram(y, 100, 'Normalization', 'pdf', 'DisplayName',
, 'Generated Samples');
grid on;
hold on;

a = 2*c;
x_pdf = linspace(-a, a, 400);
y_pdf = (1/a) * (1 - abs(x_pdf)/a);
plot(x_pdf, y_pdf, 'r-', 'LineWidth', 2.5, 'DisplayName',
, 'Theoretical PDF');

title('Symmetric Triangular Distribution');
xlabel('Random Variable Value');
ylabel('Probability Density Function (PDF)');
legend;
hold off;

```

Question: Take $10 \cdot 2^{10}$ of these triangularly distributed samples, quantize them, and estimate the SQNR empirically for $N = 3, 4, 5$ and 6 bits. Do this for σ_x varying in the range $[-50, 0]$ dBFS and in steps of $0,1$ dBFS. Plot the resulting curves (SQNR in dB vs. σ_x in dBFS) along with the theoretical expression

$$\text{SQNR} = 6,02N + 4,77 - 20 \log_{10} \frac{\text{FS}}{\sigma_x} \quad (\text{dB}). \quad (1)$$

Are there any differences between the theoretical and empirical curves? If so, how do you explain them? ADD

Question: In view of your results, what are the optimum values (regarding SQNR) of σ_x (in dBFS), and for the different resolutions analyzed (3 to 6 bits)? Does this agree with your intuition (see first point above)?

4 Task 4

Assume a full-scale sinusoidal input with $f_0 = 37,1094\text{MHz}$, and let the FFT size be $M = 1024$. Generate $15 \cdot M$ samples of $x(t)$ (at $f_s = 100\text{MHz}$) and quantize them to $N = 12$ bits. Break the vector x_q of quantized samples into 15 size- M blocks using, e.g., the command `reshape`:

```
xqblocks = reshape(xq, M, 15);
```

so that each column of the $M \times 15$ matrix `xqblocks` will contain the corresponding block of size M . Now, since the `fft` command computes the FFT columnwise, in order to apply an M -point FFT to each block, we simply make

```
X = fft(xqblocks, M);
```

Average the squared magnitude of the DFT coefficients over the 15 blocks and plot the results between 0 and $f_s/2$, in dBFS. Observe the location and peak value of the principal frequency component, as well as the value of the noise floor. Do your observations agree (quantitatively) with what you would expect?